


Lesson 10: Confidence intervals

TB sections 4.2

Meike Niederhausen and Nicky Wakim

2024-11-04

Where are we?



Learning Objectives



1. Calculate a confidence interval when we know the population standard deviation
2. Interpret a confidence interval when we know the population standard deviation
3. Calculate and interpret a confidence interval *using the t-distribution* when we do not know the population standard deviation

Learning Objectives

1. Calculate a confidence interval when we know the population standard deviation
2. Interpret a confidence interval when we know the population standard deviation
3. Calculate and interpret a confidence interval *using the t-distribution* when we do not know the population standard deviation

Last time: Central Limit Theorem applied to sampling distribution

- CLT tells us that we can model the sampling distribution of mean heights using a normal distribution



$$\bar{X} \sim \text{Normal}(\mu_{\bar{X}} = 65, SE = 0.424)$$

Last time: Sampling Distribution of Sample Means (with the CLT)

- The **sampling distribution** is the distribution of sample means calculated from repeated random samples of *the same size* from the same population
- It is useful to think of a **particular sample statistic** as being drawn from a **sampling distribution**
 - So the red sample with $\bar{x} = 65.1$ is just one sample mean in the sampling distribution


With CLT and \bar{X} as the RV for the **sampling distribution**

- **Theoretically** (using only population values):

$$\bar{X} \sim \text{Normal}(\mu_{\bar{X}} = \mu, \sigma_{\bar{X}} = SE = \frac{\sigma}{\sqrt{n}})$$

- **In real use** (using sample values for SE):


$$\bar{X} \sim \text{Normal}(\mu_{\bar{X}} = \mu, \sigma_{\bar{X}} = SE = \frac{s}{\sqrt{n}})$$



$$\mu_{\bar{X}} = 65 \text{ inches}$$

$$SE = 0.424 \text{ inches}$$

Last time: point estimates



Sample 50 people
 $\bar{x} = 65.1, s = 2.8$

Sample 50 people
 $\bar{x} = 64.7, s = 3.1$

Sample 50 people
 $\bar{x} = 64.9, s = 3.2$

Sample 50 people
 $\bar{x} = 66.1, s = 3.4$

This time: Interval estimates of population parameter


- A **point estimate** consists of a single value
- An **interval estimate** provides a **plausible range of values for a parameter**
 - Remember: parameters are from the population and estimates are from our sample
- We can create a plausible range of values for a population mean (μ) from a sample's mean \bar{x}
- A **confidence interval** gives us a plausible range for μ
- Confidence intervals take the general form:

$$(\bar{x} - m, \bar{x} + m) = \bar{x} \pm m$$


- Where m is the margin of error

Point estimates with their confidence intervals for μ


Sample 50 people
 $\bar{x} = 64.7, s = 3.1$




Sample 50 people
 $\bar{x} = 64.9, s = 3.2$



Sample 50 people
 $\bar{x} = 65.1, s = 2.8$



Sample 50 people
 $\bar{x} = 66.1, s = 3.4$



Do these confidence intervals include μ ?

Poll Everywhere Question 1

Confidence interval (CI) for the mean μ

Confidence interval for μ

$$\bar{x} \pm z^* \times \text{SE}$$

- with $\text{SE} = \frac{\sigma}{\sqrt{n}}$ if population sd is known

When can this be applied?

- When CLT can be applied!
- When we know the population standard deviation!

- z^* depends on the confidence level
- For a 95% CI, z^* is chosen such that 95% of the standard normal curve is between $-z^*$ and z^*
 - This corresponds to $z^* = 1.96$ for a 95% CI
- We can use R to calculate z^* for any desired CI
- Below is how we calculate z^* for the 95% CI

```
1 qnorm(p = 0.975)  
[1] 1.959964
```

Example: CI for mean height μ with σ

Example 1: Using our green sample from previous plots

For a random sample of 50 people, the mean height is 66.1 inches. Assume the population standard deviation is 3 inches. Find the 95% confidence interval for the population mean.

$$\bar{x} \pm z^* \times \text{SE}$$

$$\bar{x} \pm z^* \times \frac{\sigma}{\sqrt{n}}$$

$$66.1 \pm 1.96 \times \frac{3}{\sqrt{50}}$$

$$66.1 \pm 0.8315576$$

$$(66.1 - 0.8315576, 66.1 + 0.8315576)$$

$$(65.268, 66.932)$$

We are 95% confident that the mean height is between 65.268 and 66.932 inches.

Learning Objectives

1. Calculate a confidence interval when we know the population standard deviation
2. Interpret a confidence interval when we know the population standard deviation
3. Calculate and interpret a confidence interval *using the t-distribution* when we do not know the population standard deviation

How do we interpret confidence intervals? (1/2)


Simulating Confidence Intervals:

<http://www.rosmarchance.com/applets/ConfSim.html>

The figure shows CI's from 100 simulations:

- The true value of $\mu = 65$ is the vertical black line
- The horizontal lines are 95% CI's from 100 samples
 - **Blue**: the CI “captured” the true value of μ
 - **Red**: the CI *did not* “capture” the true value of μ

What percent of CI's captured the true value of μ ?



How do we interpret confidence intervals? (2/2)

Actual interpretation:

- If we were to
 - repeatedly take random samples from a population and
 - calculate a 95% CI for each random sample,
- then we would expect 95% of our CI's to contain the true population parameter μ .

What we typically write as “shorthand”:

- In general form: We are 95% confident that (the 95% confidence interval) captures the value of the population parameter.

WRONG interpretation:

- There is a 95% chance that (the 95% confidence interval) captures the value of the population parameter.
 - For one CI on its own, it either does or doesn't contain the population parameter with probability 0 or 1. We just don't know which!


Poll Everywhere Question 2

Learning Objectives

1. Calculate a confidence interval when we know the population standard deviation
2. Interpret a confidence interval when we know the population standard deviation
3. Calculate and interpret a confidence interval *using the t-distribution* when we do not know the population standard deviation

What if we don't know σ ? (1/2)

Simulating Confidence Intervals: <http://www.rosmarchance.com/applets/ConfSim.html>



- The normal distribution doesn't have a 95% "coverage rate" when using s instead of σ
- There's another distribution, called the t-distribution, that does have a 95% "coverage rate" when we use s

Poll Everywhere Question 3

What if we don't know σ ? (2/2)


- In real life, we don't know what the population sd is (σ)
- If we replace σ with s in the SE formula, we add in additional variability to the SE!

$$\frac{\sigma}{\sqrt{n}} \quad \text{vs.} \quad \frac{s}{\sqrt{n}}$$

- Thus when using s instead of σ when calculating the SE, we **need a different probability distribution** with thicker tails than the normal distribution.
 - In practice this will mean using a different value than 1.96 when calculating the CI
- Instead, we use the **Student's t-distribution**

Student's t-distribution

- Is bell shaped and symmetric
- A “generalized” version of the normal distribution
- Its tails are thicker than that of a normal distribution
 - The “thickness” depends on its **degrees of freedom**:
 $df = n-1$, where n = sample size
- As the degrees of freedom (sample size) increase,
 - the tails are less thick, and
 - the t-distribution is more like a normal distribution
 - in theory, with an infinite sample size the *t*-distribution is a normal distribution.



Confidence interval (CI) for the mean μ

Confidence interval for μ

$$\bar{x} \pm t^* \times \text{SE}$$

- with $\text{SE} = \frac{s}{\sqrt{n}}$ if population sd is not known

- t^* depends on the confidence level and degrees of freedom
 - degrees of freedom (df) is: $df = n - 1$ (n is number of observations in sample)
- `qt` gives the quartiles for a t-distribution. Need to specify
 - the percent under the curve to the left of the quartile
 - the degrees of freedom = $n - 1$
- Note in the R output to the right that t^* gets closer to 1.96 as the sample size increases

When can this be applied?

- When CLT can be applied!
- When we **do not** know the population standard deviation!

```
1 qt(p = 0.975, df=9) #df = n-1
```

```
[1] 2.262157
```

```
1 qt(p = 0.975, df=49)
```

```
[1] 2.009575
```

```
1 qt(p = 0.975, df=99)
```

```
[1] 1.984217
```

```
1 qt(p = 0.975, df=999)
```

```
[1] 1.962341
```

Example: CI for mean height μ with s

Example 2: Using our green sample from previous plots

For a random sample of 50 people, the mean height is 66.1 inches and the standard deviation is 3.5 inches. Find the 95% confidence interval for the population mean.

$$\bar{x} \pm t^* \times \text{SE}$$

$$\bar{x} \pm t^* \times \frac{s}{\sqrt{n}}$$

$$66.1 \pm 2.0096 \times \frac{3.5}{\sqrt{50}}$$

$$66.1 \pm 0.994689$$

$$(66.1 - 0.994689, 66.1 + 0.994689)$$

$$(65.105, 67.095)$$

What is t^* ?

$$df = n - 1 = 50 - 1 = 49$$

$$t^* = \text{qt}(p = 0.975, df = 49) = 2.0096$$

We are 95% confident that the mean height is between 65.105 and 67.095 inches.

Confidence interval (CI) for the mean μ (z vs. t)

- In summary, we have two cases that lead to different ways to calculate the confidence interval

Case 1: We know the population standard deviation

$$\bar{x} \pm z^* \times \text{SE}$$

- with $\text{SE} = \frac{\sigma}{\sqrt{n}}$ and σ is the population standard deviation
- For 95% CI, we use:
 - $z^* = \text{qnorm}(p = 0.975) = 1.96$

Case 2: We do not know the population sd

$$\bar{x} \pm t^* \times \text{SE}$$

- with $\text{SE} = \frac{s}{\sqrt{n}}$ and s is the sample standard deviation
- For 95% CI, we use:
 - $t^* = \text{qt}(p = 0.975, df = n-1)$

Some final words (said slightly differently?)

- Rule of thumb:
 - Use **normal distribution ONLY if you know the population standard deviation σ**
 - If **using s for the SE , then use the Student's t-distribution**
- For either case, we need to remember when we can calculate the confidence interval:
 - $n \geq 30$ and population distribution not strongly skewed (using Central Limit Theorem)
 - If there is skew or some large outliers, then $n \geq 50$ gives better estimates
 - $n < 30$ and data approximately symmetric with no large outliers
- **If do not know population distribution, then check the distribution of the data.**
 - Aka, use what we learned in datavisualization to see what the data look like

